

Creating a Prototype for a Data Anonymization Service

MASTER THESIS

Tobias Kaiser

Submitted on 30 July 2026



Friedrich-Alexander-Universität Erlangen-Nürnberg
Faculty of Engineering, Department Computer Science
Professorship for Open Source Software

Supervisor:
Thomas Wolter
Prof. Dr. Dirk Riehle, M.B.A.



Friedrich-Alexander-Universität
Faculty of Engineering

Declaration of Originality

I confirm that the submitted thesis is original work and was written by me without further assistance. Appropriate credit has been given where reference has been made to the work of others. The thesis was not examined before, nor has it been published. The submitted electronic version of the thesis matches the printed version.

Erlangen, 30 July 2026

License

This work is licensed under the Creative Commons Attribution 4.0 International license (CC BY 4.0), see <https://creativecommons.org/licenses/by/4.0/>

Erlangen, 30 July 2026

Abstract

This thesis addresses the critical conflict between data-driven software engineering analytics and individual privacy rights under the General Data Protection Regulation (GDPR). The MECOIS project transforms vast development metadata into actionable productivity insights, yet the granular nature of this information poses significant privacy risks. To resolve this, this work designs and implements a prototype data anonymization service as a decoupled, containerized system within the MECOIS architecture. The service utilizes a proxy-based approach for identity management, pseudonymizing direct identifiers into deterministic Universally Unique Identifiers (UUIDs) and achieving k-anonymity by grouping contributors into organizational equivalence classes. Evaluation of the prototype indicates that it successfully maintains data integrity for high-dimensional metadata while effectively mitigating the risk of "singling out" individuals. Ultimately, this study provides a scalable and legally grounded foundation for software analytics that respects developer privacy.

Contents

1	Introduction	1
2	Key Terms and Definitions	3
2.1	German Privacy Laws	3
2.1.1	General Data Protection Regulation	3
2.1.2	Federal Data Protection Act	3
2.2	Pseudonymization and Anonymization	4
2.2.1	Pseudonymization	4
2.2.2	Anonymization	4
2.2.3	Difference Pseudonymization and Anonymization	4
3	Literature Review	7
3.1	Data Protection Laws	7
3.2	Identity Management	8
4	Requirements	11
4.1	Functional Requirements	11
4.2	Nonfunctional Requirements	11
5	Architecture	15
5.1	Main System	15
5.1.1	Components	15
5.1.2	Data Flow	17
5.2	Anonymization System	19
5.3	Identity Management System	20
5.3.1	Identities	20
5.3.2	Front end Dashboard	20
5.3.3	API	21
6	Design and Implementation	23
6.1	Reason for an Anonymization Service	23
6.1.1	Purpose Limitation	23

6.1.2	Anonymized Data	23
6.1.3	Consequences for MECOIS	24
6.2	Identifying Identifiers	24
6.3	Approaches for Pseudonymization and Anonymization	25
6.3.1	Pseudonymization Techniques	26
6.3.2	Anonymization Techniques for Structured Data	27
6.3.3	Techniques for Unstructured Data	30
6.4	The Risk of Re-Identification	31
6.5	Deciding on an Approach	32
6.6	Implementing the Anonymization Service	32
6.6.1	Modules	32
6.6.2	Implementing the Server	34
6.6.3	Implementing the Pseudonymization	35
6.6.4	Implementing the Anonymization	35
6.6.5	Deploying the Server	35
7	Evaluation	37
7.1	Functional Requirements	37
7.2	Nonfunctional Requirements	40
8	Conclusions	43
8.1	Summary of Results	43
8.2	Discussion and Legal Compliance	44
8.3	Future Work	44
	References	49

List of Figures

5.1	Basic System Architecture	15
5.2	Main System Components (MECOIS, 2026)	16
5.3	MECOIS Frontend Data Explorer	17
5.4	MECOIS Frontend Data Extractor	18
5.5	Data flow (MECOIS, 2026)	19
5.6	SortingHat UML diagram (Dueñas et al., 2021, p. 15)	20
5.7	SortingHat front end (individuals are censored for privacy reasons)	21
5.8	SortingHat individuals section (uncensored) (CHAOSS, 2026c) . .	22
6.1	Server Diagram	34

List of Tables

4.1	Functional requirements	12
4.2	Nonfunctional requirements	13
6.1	Server Endpoints Categorized by HTTP Method	34

Acronyms

DSGVO Datenschutz-Grundverordnung

GDPR General Data Protection Regulation

FLR functional legal requirements

FSR functional system requirements

FIR functional integration requirements

NFLR nonfunctional legal requirements

NFSR nonfunctional system requirements

NFIR nonfunctional integration requirements

LLM Large Language Model

UUID Universally Unique Identifier

BDSG Bundesdatenschutzgesetz (English: Federal Data Protection Act)

ECJ European Court of Justice

GAN Generative Adversarial Network

NLP Natural Language Processing

A4NT Adversarial Author Attribute Anonymity Neural Translation

DPO Data Protection Officer

1 Introduction

Deriving actionable insights from vast datasets is fundamentally driving innovation (is a innovation driver across different industries) across diverse industries, from healthcare analytics to financial modeling. Yet, this heavy reliance on data accumulation (Yet often times, these innovations stand directly at odds with data privacy, which le presents a critical dichotomy. On one hand, organizations are legally and ethically mandated to protect individual privacy under frameworks such as the General Data Protection Regulation (GDPR) and emerging global standards. On the other, they remain wholly dependent on utilizing that same data Ich würde hier vager bleiben to train machine learning models and perform complex statistical analyses.

In the specialized field of software engineering, this conflict is centered Ich wuerde sagen, dass es in de on the vast vast ist ein Wort das die AI sehr oft verwendet, schau mal das du Synonyme findest amount of metadata generated by development teams through their daily interactions with tools like GitHub and Jira. Effectively harnessing this information to drive productivity while complying with privacy laws requires innovative frameworks that can navigate the complexities of high-dimensional data.

The MECOIS Ist das ne Abkürzung? project, a research initiative from the University of Erlangen, addresses the "Big Data" (this) challenge within the specific domain of software engineering by transforming the vast metadata generated during development into actionable auch Wortwiederholung, klingt dadurch sehr AI intelligence. Software development teams leave a continuous trail of activity-based artifacts in repositories such as GitHub and Jira erwahenst du auch zweimal identisch, brauchst creating a high-dimensional data set that is traditionally difficult to interpret (is difficult to interpret traditionally) due to overlapping metrics, complexity, and incompatible value ranges. MECOIS solves these issues by applying statistical methods— AI dash specifically the creation of a composite index—to group correlating metrics into key dimensions. This data-driven approach (attempts to provide) provides an objective Productivity Index that allows organizations to compare performance across teams and projects, identifying early-stage bottlenecks like miscommunication or resource gaps while uncovering the best practices of high-performing teams. (Riehle, 2025)

1. Introduction

To address the tension between data utilization and privacy, this thesis proposes a prototype for a dedicated anonymization service designed to process the metadata collected within the MECOIS framework. While MECOIS extracts high-dimensional, activity-based artifacts from platforms such as GitHub and Jira schon wieder to provide deep deep streichen insights into team productivity, the granular nature of this data — AI dash including timestamps, commit details, and developer identities — presents significant privacy risks. The proposed service aims to anonymize this data in strict accordance with the GDPR. Hier vielleicht noch nicht namedroppen sondern einfach sagen mit Gesetz und ethischen Massstäben ensuring that the trail of artifacts used to calculate the MECOIS Productivity Index is stripped of personally identifiable information. By integrating this anonymization layer, the project seeks to enable the MECOIS's mission of empowering software developers through data-driven analysis while maintaining a privacy-first approach that meets modern legal and ethical standards.

2 Key Terms and Definitions

2.1 German Privacy Laws

In Germany data privacy is regulated in two laws: The Datenschutz-Grundverordnung (DSGVO) and the Bundesdatenschutzgesetz (English: Federal Data Protection Act) (BDSG).

2.1.1 General Data Protection Regulation

The DSGVO (often also abbreviated to DS-GVO) is the German name for the GDPR, which is a European Union regulation designed to ensure the protection of personal data and the free flow of data in Europe. It has been mandatory in all EU member states since May 25, 2018, and applies equally to private companies, public authorities, associations, and individuals. The GDPR strengthens the rights of EU citizens, as they now enjoy the same legal certainty in all member states and can decide how their own data is processed. For companies, it simplifies cross-border activities through uniform rules of conduct and the so-called “one-stop-shop” procedure, but also provides for substantial fines in the event of violations. While it is directly applicable in all EU member states, individual countries still have their own national data protection laws to supplement it. (für den Datenschutz und die Informationsfreiheit [BfDI], 2025)

2.1.2 Federal Data Protection Act

Germany passed its own privacy laws in 1977 called the BDSG. It had to be revised upon the adoption of the GDPR. Today the BDSG and DSGVO complement each other, since the GDPR has escape clauses, where the BDSG steps in. This applies, for example, to the processing of employee data or the establishment of supervisory authorities. (BfDI, 2025) (Däubler, 2017, pp. 68–70)

2.2 Pseudonymization and Anonymization

To comply with the data privacy laws, data often has to be pseudonymized or anonymized.

2.2.1 Pseudonymization

According to Article 4 No. 5 of the GDPR, pseudonymization is defined as the processing of personal data in such a manner that it can no longer be attributed to a specific data subject without the use of additional information. This additional information, often referred to as a "key" for re-identification, must be kept separately and be subject to rigorous technical and organizational measures to ensure that the personal data is not assigned to an identified or identifiable natural person. Crucially, pseudonymization is characterized as a processing operation rather than a static state, involving the active transformation of cleartext data into pseudonyms to mitigate risk while maintaining the utility of the data set. Because the underlying connection to a natural person is merely obscured rather than permanently severed, pseudonymized data remains classified as "personal data" and thus stays fully within the legal scope and obligations of the GDPR. (Members of the "Deutsche Gesellschaft für Medizinische Informatik et al. [GMDS], 2024, pp. 8, 21, 71) (Schwartzmann et al., 2022, p. 14)

2.2.2 Anonymization

In contrast, anonymization is the process of transforming personal data into anonymous information so that the data subject is no longer identifiable. While not explicitly defined within the text of the GDPR itself, it is recognized as a process that results in data that does not relate to an identified or identifiable natural person. A successful anonymization is fundamentally irreversible; once the process is complete, a re-identification of the individual must be impossible for the controller or any third party. Within the academic and legal discourse, a distinction is often made between "absolute anonymization," where re-identification is factually impossible for everyone, and "factual" or "relative anonymization," where re-identification would require a disproportionate effort in terms of time, cost, and labor. Once data is effectively anonymized, it is no longer considered personal data, and the requirements of the GDPR cease to apply. (GMDS, b., BvD, 2024, pp. 1, 9–10, 13, 21, 68, 73) (Schwartzmann et al., 2022, pp. 3, 16–19, 23)

2.2.3 Difference Pseudonymization and Anonymization

The fundamental difference between the two terms lies in the reversibility of the process and the subsequent legal classification of the data. Pseudonymization

is a reversible procedure where the link between the data and the individual is preserved via a separate key, meaning the data retains its character as "personal data". Anonymization, however, aims for an irreversible state that permanently breaks this link, thereby removing the data from the regulatory framework of the GDPR. Furthermore, pseudonymization requires the controller to implement strict silos and access controls to manage re-identification information, whereas anonymization requires the permanent deletion or sufficient generalization of all direct and indirect identifiers to ensure that no such "key" exists to restore the original identity. (GMDS, b., BvD, 2024, pp. 8, 13, 31, 68) (Schwartmann et al., 2022, pp. 3, 19)

2. Key Terms and Definitions

3 Literature Review

This section reviews existing scientific publications to establish a baseline for data anonymization services within software engineering productivity metrics.

3.1 Data Protection Laws

While extensive literature exists on the GDPR and data anonymization, existing studies typically address these concepts at a high level or focus primarily on surveillance applications. Consequently, domain-specific literature directly aligned with the scope of this research remains scarce.

The standard reference work on employee data protection is Däubler (2017). Written in a practical style and accessible even to non-specialists, this handbook explains the entire field of data protection law. It focuses on employee data protection and the participation rights of works councils and staff councils. It also addresses the challenges posed by AI that are already apparent today, such as those arising in the hiring process or when working with ChatGPT. Another key focus is on the diverse means available to ensure that data protection is effectively enforced in practice. These range from data backup and the activities of the company data protection officer and the supervisory authority to the imposition of substantial fines and numerous claims for damages. The now very extensive case law of the European Court of Justice (ECJ) — which has set many new precedents — is analyzed in relation to these and numerous other questions. As a rule, this case law strengthens the protection of individuals, which is also of great significance for the scope of action available to works councils and staff councils. This shift has not yet been fully recognized everywhere.

In the context of data protection and research, the operationalization of de-identification procedures under the DSGVO is addressed by two central German guidelines that provide both technical and legal frameworks. Schwartzmann et al. (2022) establishes a structured procedural model that distinguishes between theoretical "absolute anonymity" and the more practically relevant "factual anonymization," where re-identification is considered impossible due to the dispro-

portionate effort required in terms of time, cost, and manpower. Complementing this, GMDS, b., BvD (2024) focuses specifically on the sensitive domain of health data under Art. 9 GDPR, highlighting the necessity of aligning national definitions with the broader European legal framework provided by Directive (EU) 2019/1024. While Schwartmann et al. (2022) categorizes technical methods—such as randomization (e.g., differential privacy) and generalization (e.g., k-anonymity)—into specific application classes like AI algorithm training or software testing, GMDS, b., BvD (2024) provides detailed risk assessments for specific attack scenarios, including “singling out,” inference, and linkage attacks. Both sources emphasize that anonymity is not a static state but a dynamic process that must be continuously evaluated against the evolving “state of science and technology”. Furthermore, they converge on the legal requirement that both anonymization and pseudonymization constitute “processing” activities under the GDPR, therefore requiring a valid legal basis and remaining subject to the controller’s accountability obligations.

The monograph Krebs and Hagenweiler (2022) provides a comprehensive examination of data protection techniques tailored for modern, data-driven environments. The authors systematically bridge the gap between legal requirements under the GDPR and BDSG and the technical implementation of protective measures. The work categorizes attributes into direct, indirect, and sensitive identifiers to address risks such as singling out, linkage, and inference. Beyond traditional methods like k-anonymity and generalization, the book explores advanced privacy-preserving paradigms including differential privacy, synthetic data generation via Generative Adversarial Networks (GANs), and federated machine learning. By addressing the specific challenges of both structured and unstructured data — such as text and multimedia — the authors offer a dual perspective on the legal and technical “building blocks” necessary for maintaining privacy in the era of high-dimensional Big Data analysis.

3.2 Identity Management

The MECOIS project uses the tool SortingHat from the tool set GrimoireLab by CHAOSS for managing identities.

Dueñas et al. (2021) present GrimoireLab, an integrated and modular open-source tool set developed to facilitate the comprehensive retrieval, storage, and visualization of data from software repositories. This framework addresses a critical gap in software analytics by providing automated support for approximately 30 different data sources, including version control systems like Git, issue trackers such as Jira and Bugzilla, and communication platforms like Slack and mailing lists. Central to its architecture is a multi-stage pipeline that employs specialized components: Perceval for uniform data retrieval via JSON documents,

Graal for detailed source code analysis, and SortingHat for advanced identity management and cross-platform contributor merging. Furthermore, GrimoireLab utilizes a distinct storage model that maintains "raw" copies of original data alongside "enriched" indexes optimized for visualization, a strategy that significantly enhances data maintainability, traceability, and research reproducibility. By providing a mature, field-tested framework for both academic research and industrial application, the authors demonstrate that GrimoireLab reduces the manual effort required for empirical software engineering studies while enabling the longitudinal tracking of developer activity and project health.

SortingHat is the specialized identity management component within the GrimoireLab framework, specifically designed to address the challenge of contributor fragmentation across diverse software repositories. Because developers frequently utilize different usernames, aliases, and email addresses across platforms such as Git, Slack, and Jira, SortingHat provides the essential infrastructure to unify these disparate identities into unique, individual profiles. The component operates by maintaining a dedicated relational database that stores identity data, affiliation tags, and the original source of each identity to ensure research traceability. To maintain high data quality, SortingHat utilizes a conservative algorithmic approach to merging—avoiding false positives caused by shared generic accounts—which is further complemented by HatStall, a web-based interface that allows researchers to manually curate identities and manage organizational affiliations. This process of identity unification is fundamental to the broader GrimoireLab pipeline; it allows the analytics engine to enrich data points with a unique contributor Universally Unique Identifier (UUID), thereby enabling the longitudinal tracking of developer trajectories and cross-platform activity metrics that would otherwise remain siloed. (Dueñas et al., 2021)

3. Literature Review

4 Requirements

The purpose of this chapter is to define the requirements necessary to facilitate a productivity benchmarking system compliant with the GDPR and BDSG. ‘ISO/IEC/IEEE International Standard - Systems and software engineering–Vocabulary’ (2017) splits requirements into functional and nonfunctional ones. Functional requirements define what a product must do and nonfunctional ones declare how the system should operate. Furthermore, the requirements can be classified into legality, system and integration. The service must abide to relevant data protection laws. System requirements specify how the service itself will be designed according to the author. The MECOIS project itself sets some integration demands in its contributing guidelines (MECOIS, 2026), that must be satisfied.

4.1 Functional Requirements

The table 4.1 specifies what the anonymization must do. They are sorted by functional legal requirements (FLR), functional system requirements (FSR) and functional integration requirements (FIR).

The Data Anonymization requirement FLR-1 is defined very vague by design. There can be many ways to anonymize data according to the DSGVO. Chapter 6 will compare and present different approaches. It stands in direct conflict with FSR-1 Preserve Data. The aim is to find a compromise between these two.

4.2 Nonfunctional Requirements

The table 4.2 specifies how the anonymization service must operate. They are sorted by nonfunctional system requirements (NFSR) and nonfunctional integration requirements (NFIR). There are no nonfunctional legal requirements (NFLR), because it is evaluated in chapter 6, how the legal requirements will be satisfied.

Table 4.1: Functional requirements

ID	Requirement	Description
FLR-1	Data Anonymization	The service must follow the German GDPR (called DSGVO) and BDSG.
FSR-1	Preserve Data	The service must preserve all available Data.
FSR-2	Independent Deployment	The service can be deployed and run independent from the rest of the ME-COIS project.
FSR-3	Deterministic vs. Non-Deterministic Pseudonym Generation	The service must support consistent UUID generation across pipeline runs so that cross-platform metrics can be linked over time without revealing identity.
FIR-1	Seamless Integration	The service must integrate into the project with as few changes as possible to existing features.

Table 4.2: Nonfunctional requirements

ID	Requirement	Description
NFSR-1	Containerized Service	The service must run in a docker container.
NFSR-2	Horizontal Scalability and Statelessness	The anonymization service server must remain stateless so multiple instances can be scaled horizontally behind a load balancer during heavy pipeline jobs.
NFIR-1	Data Integrity	The service must maintain schema consistency during the serialization and deserialization of DataFrames to ensure no data loss occurs during the HTTP transmission process.
NFIR-2	Programming Language	The Backend of the MECOIS project is written in python3. Therefore, this project must also use python3. (MECOIS, 2026)
NFIR-3	License Policy	The contributing guidelines (MECOIS, 2026) state: "You are allowed to use any package that can help you do your work, as long as it is not a copyleft or proprietary license. However, it is still a good idea to ask beforehand, so you do not put in unnecessary work."
NFIR-4	Styleguide and Clean Code	The contributing guidelines (MECOIS, 2026) state: "We follow the PEP8 standard (with minor exceptions like allowing longer line limits). You can run the custom linter using: pylint. Additionally, your code must follow these rules before being merged into the main branch: • Remove comments that are not necessary for understanding your code • Remove unnecessary print statements • Make sure to update and provide the DocString description of the methods you write. Make sure to have it adjusted to your use-case if you copy method-stubs from other part of the code"

4. Requirements

5 Architecture

The MECOIS setup consists of 3 separate systems, that communicate via HTTP (see Figure 5.1). The systems are separated to protect sensitive information.




Figure 5.1: Basic System Architecture

5.1 Main System

The main systems is responsible everything but the identity management and therefore has many components.

5.1.1 Components

Figure 5.2 displays an overview of all components. (MECOIS, 2026)

Sources

The sources define the origin for the data. Currently we are using Git (GitHub and GitLab), Confluence and Jira as source. For Example from GitHub we pull among other commits, issues and pipelines.

Infrastructure

The infrastructure provides the database for the project. The data is stored in a data lake, which is a repository, where data is stored as-is without further




Figure 5.2: Main System Components (MECOIS, 2026)

structure (Services [AWS], 2026). Each set of data passes through 4 Stages: raw, bronze, silver, gold.

Pipeline

The pipelines transform the data between the stages.

Analytics

In the last step the data is used to provide multiple analytics for customers. Examples for such analytics are:

- Calculating the Costs of Inner Source Collaboration by Computing the Time Worked (Buchner & Riehle, 2021)
- Estimating Software Contribution Time from Git and GitHub Metadata (Stenzel, 2026)

Services

Services are helper modules for pipelines and analytics. Such helper modules are:

- AI: a framework for communicating with a Large Language Model (LLM).
- config handling: Parses configuration files for pipelines.

Front end

The front end provides a Dashboard for the user-facing components of MECOIS. Figure 5.3 shows the Data Explorer tool, where users can examine the data in the data lake. Figure 5.4 presents the Data Extractor tool, that can be used to query data from the sources and feed it to the pipelines.

MECOIS

MECOISAdmin

Dashboards

Overview

Personal

Organization

Data

Pipeline

Data Explorer

Solution Demos

Community Features

Project Hub

Surveys

Settings

Logout

Search

Figure 5.3: MECOIS Frontend Data Explorer

Orchestrators

The orchestrators can be used to run the pipeline. The set up the configuration and run the pipeline for the selected sources.

5.1.2 Data Flow

After pulling the source the raw data is stored exactly as provided as JSON. In the data lake the Data is processed multiple times and the states are called bronze, silver and gold. The processing is done by the pipelines.

Each pipeline has 3 steps:

1. Read the data from the data lake
2. transform data
3. write the transformed data to the data lake






Figure 5.5: Data flow (MECOIS, 2026)

Anonymize Data The data is sent to the anonymization service for processing.

add UUID Each table has a way to create an UUID by concatenating some columns and hashing the result.

Silver to Silver

Identities may update in the identity management system. The `UpdateIdentityPipeline` updates any identities in the data lake that have been updated in the SortingHat.

Silver to Gold

There is no single `Silver2GoldPipeline`. The analytics represent the gold to silver pipelines. Each analytic uses the silver data to provide some metric.

Gold to Gold

Some analytics may need some mechanism to update the generated metric. This process is defined as gold to gold transformation.

5.2 Anonymization System

The anonymization system provides the anonymization service, that anonymizes the data according to the GDPR and BDSG. It provides a client library for sending data and requesting identities.

5.3 Identity Management System

The identity management system is responsible for storing all information about individuals. The project is using the tool SortingHat from the tool set Grimoire-Lab by CHAOSS. This application is provided under the GPL-3.0 license.

One individual may have multiple identities on different platforms. The system can pool these different accounts into one individual. These individuals can also be grouped into teams. Multiple teams can form an organization. (CHAOSS, 2026a) (CHAOSS, 2026b)

5.3.1 Identities




Figure 5.6: SortingHat UML diagram (Dueñas et al., 2021, p. 15)

Figure 5.6 describes how identities are managed in the SortingHat. One individual may have multiple identities on different platforms. The system can pool these different accounts into one individual. One individual has a UUID and is linked to a profile containing personal information. These individuals can also be grouped into organizations. (CHAOSS, 2026a) (CHAOSS, 2026b)

5.3.2 Front end Dashboard

The SortingHat provides a dashboard, where individuals can be merged and associated into teams and organizations.

Figure 5.7 shows a screenshot of the front end. Since the screenshot contains real data, the individuals are blacked out. The individuals section can be seen in figure 5.8.




Figure 5.7: SortingHat front end (individuals are censored for privacy reasons)

5.3.3 API

The SortingHat also provided an API for automating interactions, like listing, updating and deleting individuals.

5. Architecture

Individuals 9					ADD	
☐ Select all	Filters	Search	Sort: Name	MERGE	DELETE	
EM	Eva Millán	evamillanmartin@gmail.com				
	Georg J.P. Link	linkgeorg@gmail.com				
JS	John Smith	johnsmith@gmail.com				
	Quan Zhou	quan@bitergia.com				
	Santiago Dueñas	sduenas@bitergia.com				
VS	Veerasamy Sevagen amFOSS	sevagenv@gmail.com				
VS	Veerasamy Sevagen	sevagenv@gmail.com				
	Venu Vardhan Reddy Tekula	venuvardhanreddytekula8@gmail.com				
VT	Venu Vardhan Reddy Tekula	venu@bitergia.com				

Figure 5.8: SortingHat individuals section (uncensored) (CHAOSS, 2026c)

6 Design and Implementation

Chapter 5 described the architecture of the MECOIS project and the integration of the anonymization service. This chapter will explain how the service is implemented.

6.1 Reason for an Anonymization Service

According to the BDSG, personal data may only be collected, processed, or used for a specific purpose. It is forbidden to use this data for other purposes. There are only a few specific Exceptions for this rule. (Däubler, 2017, p. 72)

6.1.1 Purpose Limitation

Since software development tools like GitHub, GitLab or Jira contains personal data (e.g. git commit author name or Jira ticket reporter), its use falls under the GDPR. In most cases there is no written agreement on how the data can be used, so it can be assumed that the purpose for this data is to help developers organize their work. Generating productivity measurements for individual employees would be considered a misuse of the data. Däubler (2017, pp. 266–267) explicitly states that data that consists of operational workflows gathered by electronic systems cannot be used for measuring productivity of individual employees. (Däubler, 2017, pp. 266–267)

One exception is scientific research. If individuals are part of a scientific research project, their data can be used for this purpose. The data should get pseudonymized or anonymized as fast as possible, in order to minimize the infringement of privacy of the individual members.(Däubler, 2017, pp. 273–274)

6.1.2 Anonymized Data

The BDSG only applies to personal Data, which is defined in Section 3 (1): Personal data consists of specific information about the personal or factual circumstances of an identified or identifiable natural person (data subject). (Däubler,

2017, p. 69)

The DSGVO has an similar approach. In Article 4(1) it defines personal data as any information relating to an identified or identifiable natural person (hereinafter referred to as the “data subject”). A natural person is considered identifiable if they can be identified, directly or indirectly, in particular by association with an identifier such as a name, an identification number, location data, an online identifier, or one or more factors specific to the physical, physiological, genetic, mental, economic, cultural, or social identity of that natural person. (Däubler, 2017, p. 69)

Unlike anonymized data, where re-identification is impossible, pseudonymized data is still traceable to the original person. therefore pseudonymized data has to comply with the DSGVO and BDSG. Anonymized data is not mentioned in the DSGVO because, since it does not relate to specific individuals, it is not subject to the DSGVO. (Däubler, 2017, p. 69)

6.1.3 Consequences for MECOIS

For the MECOIS project the exception for scientific research is very important regarding the development of the project. As long as MECOIS is a scientific research project, the data can be used to generate productivity measurements as long as the process serves scientific research purposes. But the data still should be pseudonymized or anonymized.

Regarding the final product, MECOIS must anonymize the Data. Otherwise it is illegal to produce productivity metrics.

6.2 Identifying Identifiers

The identification and categorization of identifiers is a fundamental prerequisite for implementing effective data protection measures, such as pseudonymization or anonymization. Attributes within a data set are typically classified into three categories: direct identifiers, indirect (or quasi-) identifiers, and sensitive attributes. (Krebs & Hagenweiler, 2022, pp. 87–88, 131–132) (Schwartmann et al., 2022, pp. 9–10) (GMDS, b., BvD, 2024, p. 37)

Direct Identifiers Direct identifiers, such as civil names, social security numbers, or email addresses, allow for the immediate identification of a natural person without further processing. (Krebs & Hagenweiler, 2022, pp. 87–88) (GMDS, b., BvD, 2024, pp. 37, 71)

Indirect Identifiers Indirect identifiers or quasi-identifiers include data points like birth dates, ZIP codes, or rare job titles that do not identify a person in isolation but can lead to identification when combined with external datasets or background knowledge — a risk often termed "singling out". (Krebs & Hagenweiler, 2022, pp. 87–88) (Schwartmann et al., 2022, pp. 5, 13) (GMDS, b., BvD, 2024, p. 37)

Sensitive Attributes Sensitive attributes, such as the special categories of data protected under Article 9 of the GDPR (e.g., genetic or health data), require the highest level of protection due to the potential for discrimination and significant impact on individual freedoms. (Krebs & Hagenweiler, 2022, pp. 131–133) (GMDS, b., BvD, 2024, pp. 8, 24–25)

Ultimately, the classification of any specific data point is context-dependent, as even non-identifying data may become identifying depending on the "reasonable" means and background information available to a potential attacker. (GMDS, b., BvD, 2024, pp. 36–37, 54)

Example GitHub Commit In the context of GitHub commits, direct identifiers are: names and emails of author and committer. Since every employee has access to the GitHub repositories, indirect identifiers like the commit message (including the timestamp), ids and specific URLs can be easily resolved to the involved persons. Sensitive attributes are not present in GitHub.

6.3 Approaches for Pseudonymization and Anonymization

Techniques for pseudonymization and anonymization are categorized based on their technical approach and the type of data they address.

The boundary between pseudonymization and anonymization is fundamentally fluid because the technical methods used — such as masking, shuffling, or the variance method — are often identical for both processes. Technically, these randomization and transformation techniques serve as common building blocks; the legal and practical classification of the resulting data depends on the management of "additional information" and the context of the data environment. For instance, a method like masking is classified as pseudonymization if the researcher retains a secure mapping table to restore the original identifiers, but it becomes anonymization if that key is irreversibly destroyed. This status remains dynamic rather than absolute, as the availability of new external datasets or advances in analytical power can enable the re-identification of data previously considered anonymous. Consequently, the transition between these two states

depends heavily on the "reasonability test", where data is only truly anonymous if re-identification is either factually impossible or requires a disproportionate effort in terms of time, cost, and manpower. (Schwartzmann et al., 2022, pp. 3–4) (GMDS, b., BvD, 2024, pp. 44–45, 55, 68, 71) (Krebs & Hagenweiler, 2022, pp. 77–80)

6.3.1 Pseudonymization Techniques

Pseudonymization involves replacing identifiers with surrogates so that the data cannot be attributed to a specific person without additional information stored separately. (Krebs & Hagenweiler, 2022, pp. 71, 77)

Masking and Replacement Direct identifiers are replaced with constant values, characters, or placeholders (e.g., replacing "John Doe" with "Person A" or setting the birthday to the first of January). (Krebs & Hagenweiler, 2022, pp. 77–80) (GMDS, b., BvD, 2024, pp. 40, 44)

Shuffling (Mixing) Values within a column are swapped or rearranged across different records through random distribution to remove the link between attributes and specific individuals. Direct identifiers have to be pseudonymized with another technique, to prevent users from unshuffling the data, by connecting direct identifiers. (Krebs & Hagenweiler, 2022, pp. 77–80) (GMDS, b., BvD, 2024, p. 41)

Variance Method For quantitative data, numerical values (e.g. ages or dates) are altered by adding random deviations within a specified range while maintaining overall statistical distribution. (Krebs & Hagenweiler, 2022, pp. 77–80) (GMDS, b., BvD, 2024, p. 42)

Cryptographic Encryption and Hashing Cryptographic methods, specifically encryption and hashing, are essential technical measures for the pseudonymization of personal data, as they replace direct identifiers with secure surrogates. Encryption involves transforming identifiers using a secret key, making the process reversible only for authorized parties who possess that key. In contrast, hashing utilizes a one-way mathematical function to map identifiers to a fixed-length string, which is inherently designed to be non-invertible. While hashing provides a more robust barrier against re-identification than simple masking, raw hash values remain vulnerable to brute-force or rainbow table attacks if an attacker can guess the range of original inputs. To mitigate this risk, it is critical to implement a "salt," which is a random sequence of characters appended to the original data before it is hashed. The use of a salt significantly increases the entropy of the input, ensuring that the resulting hash is not easily predictable and

preventing attackers from using precomputed tables to reverse the transformation. (Krebs & Hagenweiler, 2022, pp. 77–80) (GMDS, b., BvD, 2024, pp. 13, 42–44, 84)

Tokenization Often used in the financial sector, this replaces sensitive identifiers like card IDs with randomly generated tokens. (Krebs & Hagenweiler, 2022, pp. 77–80)

6.3.2 Anonymization Techniques for Structured Data

Anonymization aims to permanently remove the connection between data and a natural person, making re-identification practically impossible. (Krebs & Hagenweiler, 2022, pp. 73–75)

Structured Data Structured data refers to information that is organized according to a fixed schema or predefined format, ensuring that its content and meaning are clearly and uniquely determinable. It is fundamentally understood through a "Key-Value" representation, where the "Key" assigns specific meaning (e.g., "First Name") and the "Value" contains the actual content (e.g., "Anne"). In practice, this data is typically stored in tabular formats like spreadsheets or relational databases, where each row represents an individual data record and each column defines a specific attribute with a fixed data type, such as numeric values, categorical lists, or qualitative ratings. This rigid organization makes structured data highly searchable and suitable for systematic processing by automated analysis tools. (Krebs & Hagenweiler, 2022, pp. 45–46) (Schwartzmann et al., 2022, pp. 11–12)

Randomization (Perturbation)

These methods distort attribute values to decrease the probability of identifying individuals from a dataset. (Krebs & Hagenweiler, 2022, pp. 83–86)

Stochastic Overlay Characteristics are altered with random variables so that the specific values are less accurate but general distribution is preserved. (Krebs & Hagenweiler, 2022, p. 84) (Schwartzmann et al., 2022, p. 31)

Swapping A special form of stochastic overlay is swapping, where attribute values are shifted between different data records, which is particularly effective when combined with the removal of quasi-identifiers. (Krebs & Hagenweiler, 2022, pp. 84–85) (Schwartzmann et al., 2022, p. 31)

Falsification (Noise Addition) Artificial statistical noise is added to some or all data points to prevent reliable estimation of an individual’s original values. (Krebs & Hagenweiler, 2022, pp. 85–86) (Schwartzmann et al., 2022, p. 36)

Microaggregation (Clustering) Data records are grouped by similarity, and their values are replaced by a representative average (e.g., mean or median) for that group. (Krebs & Hagenweiler, 2022, p. 86)

Generalization and Aggregation

Generalization and aggregation are the technical pillars for implementing formal anonymity models in structured data. Generalization involves coarsening attribute values, such as replacing a specific street address with a ZIP code or a precise birth date with a birth year, to reduce data granularity. Aggregation, often applied through micro-aggregation or clustering, groups individual records and replaces specific values with representative summary values like means or medians. While these techniques effectively reduce the risk of “singling out” individuals, they inherently involve a critical trade-off: increasing the level of protection leads to a corresponding loss of data utility for analytical purposes. (Krebs & Hagenweiler, 2022, pp. 86–89, 131, 138) (Schwartzmann et al., 2022, pp. 33–34)

k-Anonymity The k-Anonymity model provides a standard for measuring anonymity by ensuring that every record in a dataset is indistinguishable from at least k-1 other records based on its quasi-identifiers. By grouping individuals into these “equivalence classes,” the model prevents attackers from isolating a specific person within the data. While a higher value for k significantly reduces the probability of identity disclosure—calculated as $1/k$ —the model remains vulnerable to attacks that can derive sensitive attributes if all members of a group share the same characteristic. (Krebs & Hagenweiler, 2022, pp. 89–92, 138) (GMDS, b., BvD, 2024, pp. 44–46) (Schwartzmann et al., 2022, pp. 33–34)

l-Diversity To address the limitations of k-Anonymity, the l-Diversity model requires that each k-anonymous group contains at least l “well-represented” values for sensitive attributes. This diversity ensures that an attacker cannot definitively infer an individual’s sensitive information even if they know the person belongs to a specific group. Common variations include distinct, entropy, and recursive (c, l)-diversity, each offering different levels of protection against inference-based attribute disclosure. (Krebs & Hagenweiler, 2022, pp. 92–93) (GMDS, b., BvD, 2024, pp. 61–62) (Schwartzmann et al., 2022, p. 34)

t-Closeness The t-Closeness criterion further refines the privacy standard by requiring that the distribution of a sensitive attribute within any group remains

close to the distribution of that attribute in the overall dataset. By limiting the deviation (represented by the parameter t) between the conditional and unconditional distributions, this model prevents attackers from gaining significant knowledge about a person through group-level statistical anomalies. This is particularly useful for maintaining data utility while protecting against sophisticated inference attacks that l-Diversity might not fully mitigate. (Krebs & Hagenweiler, 2022, pp. 93–94) (GMDS, b., BvD, 2024, p. 62)

δ -Presence The δ -Presence model specifically addresses the risk of "membership disclosure," where an attacker attempts to determine whether a specific individual is included in a private dataset. This model typically operates by comparing a private dataset against a publicly available one that contains quasi-identifiers, such as a voter registry. By quantifying and limiting the probability (δ) that a person from the larger population can be linked to the specific subset, δ -Presence ensures that the mere presence of a person in a sensitive dataset remains protected. (Krebs & Hagenweiler, 2022, p. 94) (GMDS, b., BvD, 2024, pp. 62–63)

Advanced and KI-Based Methods

Differential Privacy Differential Privacy is a rigorous mathematical framework that protects individual privacy by adding controlled statistical noise to either the underlying data or query results. This perturbation ensures that the presence or absence of a single individual does not significantly alter the outcome of any statistical analysis. Unlike syntactic models, it provides a semantic guarantee of privacy that remains robust even if an attacker possesses extensive background knowledge. The level of protection is quantified by the parameter ϵ , where a smaller value indicates stronger privacy at the cost of reduced data utility. (Krebs & Hagenweiler, 2022, pp. 96–98) (GMDS, b., BvD, 2024, pp. 46–47)

Data Synthesis Data Synthesis involves building a statistical model of the original data set to generate entirely new, artificial data that reflects the properties of the original information without including real personal identifiers. Advanced techniques, such as a GAN, allow for the creation of realistic synthetic data points — such as images or medical records — that maintain the statistical distribution of the source material. This method is particularly valuable for training machine learning models or testing software when access to authentic personal data is legally restricted. While no direct link exists between the original and synthetic data, the synthesis model must be carefully steered to prevent the disclosure of unique patterns that could lead to re-identification. (Krebs & Hagenweiler, 2022, pp. 99–101) (Schwartzmann et al., 2022, pp. 34–35)

Semantic Anonymization Semantic Anonymization uses active ontologies and inferencing within a semantic AI-based system to transform sensitive information while preserving its analytical meaning. The process involves dividing a data set into specific "data pools" to apply separate transformation rules, a method known as semantic micro-aggregation. By focusing on the meaning and statistical relationships of data points rather than just syntactic structures, it effectively prevents re-identification through quasi-identifiers. This approach aims to maximize data utility for industrial analysis while ensuring that a person-specific trace back to the original records is impossible. (Krebs & Hagenweiler, 2022, pp. 103–107)

Federated Machine Learning Federated Machine Learning is a decentralized training approach where algorithms are trained locally on end-user devices or separate company databases rather than in a central repository. Only non-personal model parameters, such as gradients, are transmitted to a central coordinator to be aggregated into a global model. Because the raw training data never leaves its original location, this method significantly reduces the risks associated with data aggregation and unauthorized access under the GDPR. This "bringing the algorithm to the data" approach allows for the development of complex AI models across multiple partners while maintaining high data sovereignty. (Krebs & Hagenweiler, 2022, pp. 101–103) (Schwartzmann et al., 2022, pp. 7, 52–53)

6.3.3 Techniques for Unstructured Data

The following techniques for unstructured data focus on removing or altering personal identifiers in formats such as text, images, and video to prevent re-identification.

Metadata Removal This process involves deleting the decoupled metadata layer in files like PDFs or Word documents, which can contain semantically identifying fields such as author names, creation dates, and titles. It is a critical first step because metadata often exists in the surrounding system (e.g., file systems or emails) and can unintentionally reveal personal information even if the main text appears anonymous. (Krebs & Hagenweiler, 2022, pp. 107–110)

Content-Level Masking and Natural Language Processing This technique uses Natural Language Processing (NLP) to identify specific entities within a text—such as names, companies, or locations—and replaces them with placeholders or generic strings. Because manual redaction is often inefficient for large datasets, computer-linguistic methods like named entity recognition are used to automatically extract and mask these identifiers. (Krebs & Hagenweiler, 2022, pp. 107–110)

Paraphrasing and Coreference Replacement This method replaces identifying references with generic descriptions or alternative phrases based on linguistic resources like ontologies or word embeddings. By identifying "coreferences" (repeated mentions of the same person), researchers can replace specific names with broader categories (e.g., changing a specific street to "a residential area") while maintaining the semantic flow and readability of the document. (Krebs & Hagenweiler, 2022, pp. 107–112)

Author Obfuscation This technique aims to mask an individual's unique writing style, which can otherwise be identified through digital text forensics. Advanced approaches like Adversarial Author Attribute Anonymity Neural Translation (A4NT) use neural networks to "translate" a text into an identity-masking form that hides attributes like gender or age while preserving the original meaning. (Krebs & Hagenweiler, 2022, pp. 110–112)

Multimedia Processing Techniques

Multimedia processing techniques are specialized methods used to obscure personal identifiers within audio, image, and video data to prevent the re-identification of individuals by human observers or automated systems. For visual media, standard approaches include pixelation (mosaicking), which divides sensitive image areas into blocks filled with average color values, and blurring, often implemented via Gaussian filters to hide faces or license plates. More direct methods include "blanking" (cutting out regions of interest) or the use of black bars, though these can often destroy the analytical context of the background. Advanced AI-driven techniques utilize digital avatars or a GAN to replace original facial features or bodies with synthetic versions that preserve the utility of movements and gestures without revealing identity. For audio data, protection involves masking unique vocal characteristics through filtering, pitch shifting, or noise addition, though comprehensive protection must also account for identifying linguistic patterns and word choices. Ultimately, the effectiveness of these techniques depends on the strength of the transformation and the surrounding context, as re-identification remains a risk if environmental clues persist or if the distortion is insufficient to resist modern biometric algorithms. (Krebs & Hagenweiler, 2022, pp. 112–117) (GMDS, b., BvD, 2024, pp. 64–65)

6.4 The Risk of Re-Identification

There is no fixed, universally applicable threshold that, once met, automatically renders data anonymous; instead, anonymization is a contextual state determined by a "reasonability test". Because absolute anonymization — where re-identification is impossible under any circumstances — is considered practically

unachievable and is not legally required, the focus is on whether identification is unlikely or requires a disproportionate effort in terms of time, cost, and manpower. Consequently, rather than relying on a static limit, researchers must perform a comprehensive risk assessment that considers the specific data environment, the intended use-case (such as internal analysis versus public release), and the likely means available to a potential attacker. This evaluation should be formalized in a risk report or "attacker model" that documents the chosen transformation methods, identifies potential re-identification scenarios like "singling out" or "linkage attacks," and quantifies the remaining residual risk. Under the accountability principle of the GDPR, this systematic process of evaluation and documentation is essential to demonstrate that the risk to the rights and freedoms of individuals has been effectively mitigated. (Krebs & Hagenweiler, 2022, pp. 120, 124–132) (GMDS, b., BvD, 2024, pp. 18–19, 55, 60–61, 68–71) (Schwartmann et al., 2022, pp. 5–6, 17, 25–26, 28–29, 60–61)

6.5 Deciding on an Approach

Direct identifiers should be pseudonymized by hashing them into deterministic UUIDs. To achieve k-anonymity, data points can be aggregated according to the organizational hierarchy, specifically by grouping employees into their respective teams. However, indirect identifiers, such as Git commit messages, still present a re-identification risk depending on the specific metrics being analyzed; in a worst-case scenario, these messages must be completely removed. Finally, re-identification risks associated with timestamps can be mitigated by injecting temporal noise.

6.6 Implementing the Anonymization Service

To isolate the handling of sensitive data, the anonymization service operates independently from the primary MECOIS architecture. In addition to performing anonymization, the service acts as a proxy for requests directed to SortingHat, granting fine-grained control over the information exposed to MECOIS. Furthermore, this decoupled architecture supports flexible deployment models: individual components can be deployed independently, allowing MECOIS to run as a standalone application or as a fully integrated on-premise installation.

6.6.1 Modules

The anonymization service is structured into five core modules:

client.requestor A client library that provides high-level functions to issue HTTP requests to the anonymization server and process incoming responses.

server.anonymizer Responsible for content transformation. It either processes payloads directly according to configured anonymization rules or delegates identity-related queries to SortingHat.

server.server Encapsulates base server operations, including request initialization, body parsing, and route handling. Upon successful validation of incoming payloads, it delegates processing tasks to the **server.anonymizer** module.

utils.converter Since MECOIS relies on PySpark DataFrames for data processing (The Apache Software Foundation, 2026), DataFrames and their schemas must be serialized to JSON for HTTP transmission. This module provides utility functions for serialization and deserialization.

utils.sortinghat_requestor A specialized client interface designed to communicate directly with the SortingHat API.

6.6.2 Implementing the Server

To minimize external software dependencies, the implementation utilizes Python’s built-in `http.server` module (Python Software Foundation, 2026b). Table 6.1 outlines the endpoints exposed by the server, categorized by HTTP request method. Figure 6.1 represents, how the server handles each request.

Table 6.1: Server Endpoints Categorized by HTTP Method

HTTP Method	Endpoints
GET	<code>docker-compose-healthcheck</code> <code>ping</code> <code>show</code> <code>show-profile</code> <code>find-profile</code> <code>find-profiles-by-organization</code>
POST	<code>change_identities</code> <code>update_identities</code>




Figure 6.1: Server Diagram

`docker-compose-healthcheck` Provides a health-check endpoint queried periodically by Docker Compose during containerized deployment (see Section 6.6.5).

`ping` Offers a basic connectivity check to verify server availability.

`show` Proxies requests to SortingHat to retrieve the most recent UUID corresponding to a specified identifier.

`show-profile` Forwards requests to SortingHat to fetch full profile metadata associated with a given UUID.

find-profile Queries SortingHat for user profiles matching a given search term.

find-profiles-by-organization Queries SortingHat for user profiles associated with a specified organization.

change_identities Executes pseudonymization and anonymization on the data set (see Sections 6.6.3 and 6.6.4). This endpoint is invoked directly during the Bronze-to-Silver pipeline execution (see Section 5.1.2).

update_identities Re-evaluates and updates previously anonymized and pseudonymized datasets in response to upstream modifications within SortingHat. This endpoint is triggered during Silver-to-Silver pipeline operations (see Section 5.1.2).

6.6.3 Implementing the Pseudonymization

Identity management is handled by SortingHat, which assigns a UUID to each identity. The anonymization service subsequently replaces all direct identifiers — specified via a configuration file — with this assigned UUID.

SortingHat generates the UUID by hashing the user’s name, email address, and username using Python’s `hashlib.sha1` module (CHAOSS, 2026b). However, this approach presents cryptographic vulnerabilities: SHA-1 is susceptible to hash collisions, lacks resistance to brute-force attacks, and does not support salting (Python Software Foundation, 2026a). Despite these security limitations, SortingHat’s native UUID mechanism is retained during the current development phase due to its implementation simplicity, with a cryptographically secure scheme planned for future iterations.

6.6.4 Implementing the Anonymization

SortingHat supports associating individuals with specific enrollments, where an enrollment is defined by a start timestamp, an end timestamp, and an organization. This capability is leveraged to enrich the DataFrame with an additional column representing each individual’s organizational affiliation.

6.6.5 Deploying the Server

For the deployment of the server, this project utilizes Docker Compose, a tool specifically designed for managing multi-container applications through a unified YAML configuration file. This approach offers significant advantages, most notably the ability to link multiple independent microservices and initiate them

simultaneously with a single command, providing a higher-level abstraction than individual run commands. Furthermore, the use of Docker ensures consistent environments and application isolation, which prevents dependency conflicts and eases the transition from development to production. The data and application logic are inherently more secure in this setup because they are encapsulated within a container — an isolated instance that utilizes namespaces and its own local file system to remain distinct from the host machine and other services. (Jangla, 2018, pp. 3–5, 13, 28–29, 34–36, 77–81)

7 Evaluation

Following the design and implementation of the anonymization service prototype, this chapter provides a systematic evaluation of the system against the functional and nonfunctional requirements established in Chapter 4. Ultimately, this evaluation determines the extent to which the service provides a legally and technically sound basis for measuring developer productivity in a research or production environment.

7.1 Functional Requirements

FLR-1: Data Anonymization

Evaluating the quality of an anonymization process requires a multifaceted analysis that balances technical robustness, legal compliance, and data utility. The "goodness" of the resulting dataset is primarily measured through formal anonymity models — such as k -anonymity — which provide mathematical metrics for the degree of protection against identity and attribute disclosure. Beyond these syntactic criteria, the effectiveness of the transformation must be validated against the three fundamental risk scenarios: singling out (isolating an individual), linkability (connecting records across datasets), and inference (deducing sensitive values from other attributes). A robust evaluation further involves the implementation of an attacker model, which simulates realistic adversaries with varying levels of background knowledge and resources to quantify the remaining residual risk. From a legal perspective, the anonymization is considered successful if it achieves factual anonymity, meaning that re-identification is either impossible or requires a disproportionate effort in terms of time, cost, and manpower. Finally, the quality of the process is intrinsically linked to data utility; a high-quality anonymization must minimize information loss to ensure that the dataset remains suitable for the intended analytical purpose while maintaining individual privacy. (Krebs & Hagenweiler, 2022, pp. 73–75, 81–82, 86–88, 103–106, 120–121) (GMDS, b., BvD, 2024, pp. 61–63) (Schwartmann et al., 2022, pp. 3–6, 28–29)

A comprehensive quantitative evaluation of every potential re-identification scenario or the implementation of advanced privacy models such as l -diversity and t -closeness is beyond the scope of this thesis. Instead, the prototype focuses on establishing a robust baseline for privacy by first pseudonymizing direct identifiers, which is achieved by hashing them into deterministic UUIDs within the identity management system. Building upon this layer of protection, k -anonymity is achieved through the aggregation of data points based on organizational hierarchy. This is specifically realized by grouping individuals by their enrollment, a mechanism that links contributors to specific organizations over defined time intervals. By organizing data into these equivalence classes, the service ensures that each record becomes indistinguishable from at least $k-1$ other individuals, thereby effectively mitigating the risk of singling out specific developers while preserving the utility of group-level productivity metrics.

FSR-1: Preserve Data

The evaluation of FSR-1: Preserve Data reveals an inherent technical and legal conflict with the requirement for data anonymization (FLR-1), as the absolute preservation of all source data is impossible when specific attributes must be transformed or removed to protect individual privacy. In the current prototype, this requirement is addressed through a tiered approach: direct identifiers are pseudonymized by hashing them into deterministic UUIDs, which remain resolvable through the SortingHat identity management system to maintain longitudinal traceability. To further protect privacy, individual contributors are generalized into equivalence classes based on their organizational enrollment, effectively achieving k -anonymity. While the retention of pseudonymized data is permissible and necessary during the scientific research and development phase of MECOIS, it may be mandatory to irreversibly delete the mapping "keys" in the final product to ensure compliance with the GDPR's anonymization standards. Consequently, the preservation of data utility remains dynamic; indirect identifiers, such as Git commit messages, may also be subject to future deletion if they pose a disproportionate re-identification risk.

FSR-2: Independent Deployment

The prototype successfully achieves architectural and operational decoupling from the primary MECOIS framework. By design, the anonymization service operates as an isolated system that communicates via standard HTTP interfaces, allowing it to be deployed as a standalone component or as part of a fully integrated on-premise installation. From a security perspective, enclosing sensitive data within these independent systems establishes a strict isolation boundary, thereby reducing the risk of unauthorized data exposure and preventing sensitive payloads from leaking into neighboring host components or the primary analytics engine.

This technical independence is primarily realized through the use of Docker Compose, which encapsulates the service and its specific environment within isolated containers, thereby preventing dependency conflicts and ensuring consistent behavior across different deployment environments. While the service functions as a proxy for the SortingHat, this decoupled architecture provides the necessary flexibility to scale the anonymization layer independently of the main analytics engine to meet varying performance demands.

FSR-3: Deterministic vs. Non-Deterministic Pseudonym Generation

The prototype successfully prioritizes a deterministic approach to ensure the consistency of longitudinal data analysis. By utilizing the native UUID generation mechanism within the SortingHat identity management system — which hashes direct identifiers such as names, emails, and usernames — the service ensures that the same individual is assigned a consistent pseudonym across different pipeline runs. This determinism is essential for the MECOIS framework’s goal of linking cross-platform metrics from disparate sources like GitHub and Jira without revealing the developers’ actual identities. However, the implementation is currently constrained by the use of the SHA-1 hashing algorithm, which presents cryptographic vulnerabilities such as a lack of resistance to brute-force attacks and an absence of salting. While this simplified method is functionally sufficient for the prototype’s current research phase and supports the necessary traceability, the evaluation highlights that a more secure, salted hashing scheme will be required for the final product to more effectively mitigate re-identification risks.

FIR-1: Seamless Integration

The anonymization service successfully integrates into the MECOIS framework by replacing legacy processing steps with calls to the new server. Specifically, the previous, insufficient approach to anonymization by pseudonymizing direct identifiers in the Bronze-to-Silver pipeline has been replaced by invoking the `change_identities` endpoint, ensuring that the data is anonymized as an integral part of the data flow. Similarly, the Silver-to-Silver pipeline has been adapted to utilize the `update_identities` endpoint to maintain data consistency when changes occur within the identity management system. Furthermore, existing communication with the SortingHat system has been successfully redirected through the new service, which now acts as a proxy for requests such as `show-profile` or `find-profile`. This transition fulfills the requirement for minimal disruption to existing features while providing the system with fine-grained control over the personal information exposed to the analytics engine.

7.2 Nonfunctional Requirements

NFSR-1: Containerized Service

The requirement is fully met through the implementation of `Docker Compose`, which manages the service as part of a multi-container application via a unified YAML configuration. This containerized approach ensures environmental consistency and application isolation, effectively preventing dependency conflicts that could arise during the transition from development to production. Furthermore, the encapsulation of the application logic within a container provides a secure, isolated instance that utilizes its own namespaces and local file system to remain distinct from the host machine and other integrated services. By leveraging these containerization technologies, the prototype achieves a high-level abstraction that allows for the simultaneous initiation of all necessary microservices with a single command, fulfilling the demand for a portable and robust deployment model.

NFSR-2: Horizontal Scalability and Statelessness

The anonymization service is designed as a stateless component, fulfilling the architectural requirement for efficient horizontal scaling. By functioning primarily as a transformation engine for incoming `DataFrames` and acting as a proxy for the `SortingHat` API, the service avoids maintaining local session state between individual HTTP requests. This stateless architecture ensures that multiple instances of the service can be seamlessly deployed behind a load balancer to distribute the computational load during heavy pipeline jobs without risk of data inconsistency. Furthermore, the implementation of containerization via `Docker Compose` facilitates the rapid initiation of additional isolated instances, providing the necessary elasticity to handle varying processing demands within the MECOIS framework. This decoupled design allows the anonymization layer to scale its resources independently of the main analytics engine, thereby ensuring both performance and high availability.

NFIR-1: Data Integrity

The anonymization service effectively maintains the consistency of data structures during its integration with the broader MECOIS framework. This requirement is successfully addressed by the `utils.converter` module, which manages the serialization and deserialization of `PySpark DataFrames` and their corresponding schemas into `JSON` format for HTTP transmission. By preserving schema integrity, the system ensures that no data loss occurs during the communication between the decentralized system components, thereby maintaining the quality of the activity-based artifacts required for statistical analysis. This technical measure is critical for the reliable derivation of the Productivity Index, as it ensures

that the high-dimensional metadata remains structurally sound and analytically useful throughout the entire anonymization pipeline.

NFIR-2: Programming Language

The requirement is fully satisfied, as the anonymization service is implemented entirely in `Python3`. This alignment with the existing MECOIS backend ensures seamless interoperability between the system components, particularly for the processing of `PySpark DataFrames` and the integration of the service into the multi-stage data pipeline. Furthermore, the use of `Python` allows the prototype to leverage standardized libraries, such as `http.server` for the backend implementation and `hashlib` for generating deterministic UUIDs, while adhering to the project's mandated coding and license policies. By fulfilling this requirement, the service maintains technical consistency across the research framework, facilitating easier maintenance and future development.

NFIR-3: License Policy

The anonymization service adheres to the project's mandate to avoid the use of packages with copyleft or proprietary licenses. The prototype is implemented using `Python` and relies primarily on its standard libraries, such as `http.server` for the backend, which conform to permissive licensing standards. Furthermore, while the service interfaces with the SortingHat identity management system—which is governed by a `GPL-3.0` license—it maintains a strictly decoupled relationship by communicating exclusively via `HTTP`. This architectural separation ensures that the service's internal codebase remains independent of external copyleft requirements, thereby fulfilling the integration constraints established by the MECOIS contributing guidelines.

NFIR-4: Styleguide and Clean Code

The implementation of the anonymization service strictly adheres to the MECOIS contributing guidelines to ensure high software quality and long-term maintainability. The codebase is developed in alignment with the `PEP8` standard. By fulfilling these requirements, the service ensures technical consistency with the rest of the MECOIS backend, facilitating a seamless review and merging process into the main branch.

8 Conclusions

This thesis has presented the design, implementation, and evaluation of a prototype data anonymization service for the MECOIS framework, addressing the critical tension In Einleitung ok, hier too much between software engineering productivity analytics and individual privacy rights under the GDPR and BDSG. By integrating this service as a decoupled layer within the MECOIS architecture, the project demonstrates a viable path for organizations to derive actionable insights from developer metadata while maintaining a privacy-first approach.

8.1 Summary of Results

The developed prototype successfully fulfills the core functional and nonfunctional requirements established for a GDPR and BDSG-compliant benchmarking system. Key achievements of this work include:

Architectural Independence The service operates as a stateless, containerized system using Docker Compose ich würde das abstrakter beschreiben als mit dem Produktnamen allowing for independent deployment and horizontal scalability without interfering with the primary MECOIS analytics engine. Außerdem Wiederverwendbarkeit oder Skalierbarkeit

Privacy Mechanisms The service establishes a robust baseline for privacy by pseudonymizing direct identifiers into deterministic UUIDs and achieving k-anonymity through the aggregation of contributors into organizational equivalence classes.

Seamless Integration By acting as a proxy for the SortingHat identity management system and providing dedicated endpoints for identity transformation, the service integrates into the existing Bronze-to-Silver and Silver-to-Silver pipelines with minimal disruption.

Data Integrity Through the `utils.converter` module, the service ensures that PySpark DataFrames and their schemas are preserved during HTTP trans-

mission, maintaining the analytical utility of high-dimensional metadata. Ich glaube die Vorteile v

8.2 Discussion and Legal Compliance

From a legal perspective, the prototype bridges the gap between scientific research and final product requirements. Der Satz funktioniert nicht richtig, scientific research bedeutet ni
While the current use of pseudonymization is sufficient for the research phase of MECOIS, the service provides the necessary (foundation to build an) infrastructure to transition toward factual anonymization — AI dash where re-identification requires disproportionate effort — by generalized grouping and the potential permanent deletion of re-identification keys. The evaluation confirms that the system effectively mitigates the risk of "singling out" individuals through its implementation of equivalence classes based on organizational enrollment.

8.3 Future Work

Despite the successful implementation of the prototype, several areas for future development remain to enhance the service's security, privacy robustness, and production readiness.

Assuring k-anonymity

While the current prototype achieves a baseline for k-anonymity by grouping contributors into organizational equivalence classes, it lacks an automated mechanism to enforce the k threshold. Future iterations must implement a check to ensure that every organization in a dataset contains at least k enrolled individuals; Warum kein Punkt? if this threshold is not met, the system should automatically aggregate smaller organizations into broader clusters to maintain the necessary anonymity level. Furthermore, the system must account for temporal risks such as employee vacations; if a query targets a narrow timeframe where several contributors are inactive, the number of actual records may fall below the k limit even if the total enrollment is sufficient. Implementing dynamic checks for active record counts within queried intervals is essential to prevent re-identification through "singling out".

Permission System and Logging

The current service acts as a proxy for the SortingHat identity management system, but it lacks a formal access control layer. Currently, any component within the network can query the system to resolve identities. To align with GDPR accountability standards, a robust permission system and authentication layer must be implemented to restrict who can resolve UUIDs or access sensitive

profile metadata. Additionally, every query to the identity management system should be logged to provide a transparent audit trail of how and when personal data is processed.

Production-Grade Server

The prototype utilizes Python’s built-in `http.server` module to minimize dependencies during development. However, this module is not recommended for production environments as it only implements basic security checks and lacks the performance capabilities for high-volume data pipelines. Future versions should replace this with a production-grade framework such as Flask combined with Gunicorn. This upgrade would provide enhanced security, support for multiple threads, and asynchronous execution, ensuring the service can scale horizontally to handle heavy processing loads behind a load balancer. (Python Software Foundation, 2026b)

Anonymization Analysis

A critical next step is a formal, comprehensive analysis of the anonymization effectiveness. While the prototype implements key privacy strategies, ich würde hier auch noch erwähnen a complete report is needed to document the “reasonability” of the approach — AI dash evaluating whether re-identification requires a disproportionate effort in terms of time and cost. This analysis should be formalized as an “attacker model” that quantifies residual risks such as linkage or inference attacks based on the specific developer metadata processed by MECOIS.

Cryptographic Security

The service currently relies on SortingHat’s native UUID generation, which uses a SHA-1 hashing mechanism. As noted in the evaluation, SHA-1 is susceptible to hash collisions and lacks resistance to brute-force or rainbow table attacks. To improve cryptographic security, future iterations must implement a more secure, salted hashing scheme. Appending a “salt” — a random sequence of characters — before hashing will significantly increase the entropy of the identifiers and prevent attackers from using precomputed tables to reverse the transformation. eventuell Quelle? Und wenn du das schreibst kannst du auch noch SHA256 schreiben

Unstructured Data Processing

While the service successfully handles structured identifiers, indirect identifiers within unstructured data remains a significant challenge. Specifically, Git commit messages often contain timestamps, specific URLs, or linguistic patterns that could lead to re-identification. Future work should integrate advanced NLP for

content-level masking or author obfuscation techniques, such as A4NT, to mask unique writing styles while preserving the analytical utility of the messages.

Right to Erasure and Deletion Cascade

To ensure full compliance with Art. 17 GDPR, commonly referred to as the “Right to be Forgotten,” the service must implement a mechanism capable of executing a deletion cascade upon a contributor’s request. This process requires the system to irreversibly purge or “hard-anonymize” all historical mappings, pseudonym keys, and logs that allow for the re-identification of that specific individual. While the legal equivalence of anonymization and deletion is a subject of academic debate, a robust transformation that renders re-identification impossible—or possible only through a disproportionate effort in time and cost — can satisfy the requirement to make personal data unrecognizable. Within the SortingHat or similar identity management systems, this necessitates the permanent destruction of deterministic UUID mapping tables to break the link between the pseudonym and the original direct identifiers. Furthermore, achieving absolute anonymity requires that no party, including the original controller, retains the means to resolve these identifiers. Transitioning the research prototype into a legally compliant product thus (therefore) depends on the implementation of these irreversible erasure protocols to uphold the accountability principle and protect the data subject’s rights. (Krebs & Hagenweiler, 2022, pp. 71–72, 126–130) (GMDS, b., BvD, 2024, pp. 13, 18–20, 73–75) (Schwartmann et al., 2022, pp. 13–17)

drei Quellen am Ende vom Textblock würde ich nur machen wenns absolut nicht anders geht

Employee Representation

Various supervisory bodies exist to safeguard and defend employee rights. Weiterer Satz der sagt d

Data Protection Officer Under certain statutory conditions, organizations are obligated to designate a Data Protection Officer (DPO). The DPO’s responsibilities encompass providing advisory support, monitoring compliance, and liaising with the supervisory authority. Furthermore, the DPO must inform and advise controllers of their legal obligations under data protection law and audit compliance with statutory mandates, such as the GDPR. Consequently, the DPO must be integrated at an early stage during the planning and design of future digital forms of work (Däubler, 2017, pp. 392–394, 400–402).

Works Council Future implementation of MECOIS in a corporate environment requires the mandatory and active involvement of the works council (*Betriebsrat*), as productivity benchmarking directly affects employee interests. Under German law, specifically § 87 Abs. 1 Nr. 6 BetrVG, the works council possesses an enforceable right of co-determination regarding the introduction and use of

technical facilities designed to monitor the behavior or performance of employees. Since MECOIS processes granular activity-based artifacts such as timestamps and commit details, it falls under this provision regardless of whether the monitoring is the primary intent. (Däubler, 2017, p. 429) (Niedenhoff, 2024, pp. 2–6, 79–81)

The employer is legally obligated to inform the works council timely and comprehensively under § 80 Abs.2 BetrVG and must provide all necessary documentation to allow the council to evaluate the impact on the workforce. Furthermore, if the system utilizes AI for its statistical calculations or composite index generation, the works council has an explicit claim to involve external experts to assess the technical implications under §80 Abs. 3 BetrVG. To facilitate a legally compliant deployment, a shop agreement (*Betriebsvereinbarung*) must be concluded, which defines the specific boundaries of data usage. The technical measures implemented in this prototype — specifically pseudonymization via deterministic UUIDs and the achievement of k-anonymity through organizational grouping—serve as a critical baseline for these negotiations to ensure that the pursuit of team productivity does not lead to the creation of "glass employees". (does not come at the expense of workers relations) (Niedenhoff, 2024, pp. 2–4, 61–62, 72–73, 96–97, 99–101, 134, 348–351) Der Employee Representation part ist sehr lang und pas

In conclusion, the MECOIS anonymization service provides a scalable and legally grounded foundation for modern software analytics, ensuring that the pursuit of team productivity does not come at the expense of developer privacy. Der Absatz und danach dieser kurze Textstummel danach ist kein guter Abschluss, das ist zu plötz

8. Conclusions

References

- Buchner, S., & Riehle, D. (2021). Calculating the costs of inner source collaboration by computing the time worked. https://dirkriehle.com/wp-content/uploads/2021/09/Buchner_Riehle_LOCAL-VERSION__Computing_the_Costs_of_Inner_Source_Collaboration.pdf
- CHAOSS. (2026a). *Sortinghat*. GrimoireLab. Retrieved July 14, 2026, from <https://chaoss.github.io/grimoirelab-tutorial/sortinghat>
- CHAOSS. (2026b). *Sortinghat github repository*. GrimoireLab. Retrieved July 14, 2026, from <https://github.com/chaoss/grimoirelab-sortinghat>
- CHAOSS. (2026c). *Sortinghat's new interface*. GrimoireLab. Retrieved July 14, 2026, from <https://chaoss.github.io/grimoirelab-tutorial/docs/sortinghat/interface/>
- Däubler, W. (2017). *Gläserne belegschaften: Das handbuch zum beschäftigend-atenschutz* (7., a thoroughly revised and updated edition). Bund-Verlag.
- Dueñas, S., Cosentino, V., Gonzalez-Barahona, J. M., del Castillo San Felix, A., Izquierdo-Cortazar, D., Cañas-Díaz, L., & Pérez García-Plaza, A. (2021). Grimoirelab: A toolset for software development analytics. *PeerJ Computer Science*, 7, e601. <https://doi.org/10.7717/peerj-cs.601>
- für den Datenschutz und die Informationsfreiheit, D. B. (2025). *Die europäische datenschutz-grundverordnung (dsgvo)*. Retrieved July 19, 2026, from <https://www.bfdi.bund.de/DE/BfDI/Inhalte/Datenschutzpfad/DSGVO.html>
- Iso/iec/ieee international standard - systems and software engineering–vocabulary. (2017). *ISO/IEC/IEEE 24765:2017(E)*, 1–541. <https://doi.org/10.1109/IEEESTD.2017.8016712>
- Jangla, K. (2018). *Accelerating development velocity using docker: Docker across microservices* [Library of Congress Control Number: 2018962734]. Apress. <https://doi.org/10.1007/978-1-4842-3936-0>
- Krebs, H.-A., & Hagenweiler, P. (2022). *Datenanonymisierung im kontext von künstlicher intelligenz und big data: Grundlagen – elementare techniken – anwendung*. Springer Vieweg. <https://doi.org/10.1007/978-3-658-37588-1>
- MECOIS. (2026). *Mecois github repository*. MECOIS. Retrieved July 8, 2026, from <https://github.com/Mecois/mecois>

- Members of the "Deutsche Gesellschaft für Medizinische Informatik, " d. D. D. e. V. (, Biometrie und Epidemiologie e. V." (GMDS) research group "Datenschutz und IT-Sicherheit im Gesundheitswesen", e. V." (bvitg) research group "Datenschutz, " G.-I., & IT-Sicherheit". (2024, January 27). *Praxishilfe zur anonymisierung/pseudonymisierung* (Praxishilfe). Berlin/Bonn.
- Niederhoff, H.-U. (2024). *Betriebsräte in deutschland: Vom arbeiterrausschuss zum mitgestalter des arbeitslebens*. Springer Gabler. <https://doi.org/10.1007/978-3-658-44226-2>
- Python Software Foundation. (2026a). *Hashlib — secure hashes and message digests: Python 3.14.6 documentation* (3.14.6). Python Software Foundation. Retrieved July 22, 2026, from <https://docs.python.org/3/library/hashlib.html>
- Python Software Foundation. (2026b). *Http.server — http servers: Python 3.14.6 documentation* (3.14.6). Python Software Foundation. Retrieved July 23, 2026, from <https://docs.python.org/3/library/http.server.html>
- Riehle, P. D. D. (2025). *Mecois - productivity measurement*. University of Erlangen, Professorship for Open-Source Software. Retrieved July 19, 2026, from <https://www.mecois.com/solutions/productivity-measurement>
- Schwartmann, R., Jaspers, A., Lepperhoff, N., Weiß, S., & Meier, M. (2022, December). *Praxisleitfaden zum anonymisieren personenbezogener daten: Anforderungen, einsatzklassen und vorgehensmodell* (Praxisleitfaden). Stiftung Datenschutz. Berlin.
- Services, A. W. (2026). *What is a data lake?* Amazon Web Services. Retrieved July 9, 2026, from <https://aws.amazon.com/what-is/data-lake/>
- Stenzel, M. (2026). Estimating software contribution time from git and github metadata. https://oss.cs.fau.de/wp-content/uploads/2026/04/Stenzel_2026.pdf
- The Apache Software Foundation. (2026). *Quickstart: Dataframe* (4.2.0). The Apache Software Foundation. Retrieved July 23, 2026, from https://spark.apache.org/docs/latest/api/python/getting_started/quickstart_df.html